Guorui Xiao

Mobile: 609-373-5351
Email: grxiao@cs.ucla.edu
Website: xertxiao.github.io

Research Interest

I am interested in database, datastream, and machine learning systems, with the ultimate goal of building scalable data-intensive systems.

Education

- **University of California, Los Angeles** — Los Angeles, CA, USA
  Masters of Science - Computer Science; GPA: 4.0/4.0 — *Expected Graduation: Mar. 2023*
  Advisor: Carlo Zaniolo
- **University of California, Los Angeles** — Los Angeles, CA, USA
  Bachelor of Science - Computer Science; GPA: 3.77/4.0; Cum Laude — *Graduated: Dec. 2020*

Publications & Manuscripts

[P1] **A Datalog based Query Language for Supporting Recursive Query Processing over Data Streams**
Guorui Xiao, Jin Wang, Jiacheng Wu, Carlo Zaniolo. Under review by IEEE International Conference on Data Engineering (**ICDE**) 2023.

[P2] **Highly Efficient String Similarity Search and Join over Compressed Indexes**
Guorui Xiao, Jin Wang, Chunbin Lin, Carlo Zaniolo. IEEE International Conference on Data Engineering (**ICDE**) 2022, pages: 232-244.

[P3] **Demonstration of LogicLib: An Expressive Multi-Language Interface over Scalable Datalog System**
Mingda Li, Jin Wang, **Guorui Xiao**, Youfu Li, Carlo Zaniolo. ACM International Conference on Information and Knowledge Management (**CIKM**) 2022, pages: 4917–4920. (demo paper)

[P4] **Scaling state vector sync**
Varun Patil, Sichen Song, **Guorui Xiao**, Lixia Zhang. ACM Conference on Information-Centric Networking. (**ICN**) 2022, pages: 168–170 (poster paper)

[P5] **RaSQL: A Powerful Language and its System for Big Data Applications**
Jin Wang, **Guorui Xiao**, Jiaqi Gu, Jiacheng Wu, Carlo Zaniolo. ACM International Conference on Management of Data (**SIGMOD**) 2020, pages: 2673-2676. (demo paper)

[M1] **ReLiShare: Reliable Leaker Identification in Sensitive Dataset Sharing**
Zhiyi Zhang, **Guorui Xiao**, Xinyu Ma, and Lixia Zhang.

[M2] **SoK: Revealing the Architectural Design Patterns in DDoS Defense Design Space**
Zhiyi Zhang, **Guorui Xiao**, Sichen Song, Angelos Stavrou, Eric Osterweil, and Lixia Zhang.

Selected Research Projects

- **Scalable Analytics Institute (ScAi)** — University of California, Los Angeles
  Research Intern — *Dec. 2019 - Now*
  - **Streaming Data Processing System that Supports Recursive Queries [P1]**
    - Proposed a high-level query language based on Datalog for data streams to support expressing recursive queries.
    - Devised a lightweight structure *Queue-Based Index* to avoid redundant computation and further proposed an efficient query evaluation method based on it.
    - Designed and implemented a prototype datastream system (∼15k lines of codes) to verify the effectiveness of the designs.
    - Conducted experiments that showed we improved ∼10X in throughput and ∼5X in tail latency on average.
  - **Unified Compression Framework to Support String Similarity Queries [P2]**
    - Proposed the first unified framework for offline and online construction of compressed inverted index to support String Similarity Search/Join applications to avoid expensive disk I/O costs.
    - Devised algorithms to achieve near-optimal compression ratio in an online manner with tools like Kernel Density Estimation.
    - Conducted experiments that showed we improved ∼5X in memory consumption.
  - **Demonstration of RaSQL [P5]**
    - Completed a demo to demonstrate that complex queries can be expressed with RaSQL and presented a user-friendly interface to interact with the RaSQL system and monitor the query results.
    - Implemented a front end over Flask with HTML/CSS/JS, connected the front end with the RaSQL system with Py4J, prepared example queries and datasets, and contributed to the paper writing.
- **Internet Research Laboratory (IRL)** — University of California, Los Angeles
  Research Intern — *Jun. 2020 - Sep. 2020*
  - **Reliable Leaker identification via shared dataset [M1]**

    - Built a prototype system focusing on Oblivious Transfer based end-to-end sharing that realizes reliable leaker identification and Merkle Tree based credential to record the resulting shared dataset.
    - Implemented a Generative Adversarial Network (GAN) based synthetic tabular data generator on minimizing the impact on authentic shared data.
    - Prepared dataset and conducted experiments to show we achieved $< 1 \times 10^{-8}$ false negative rates by inserting only a few rows of synthetic data.
  - **Systematization of Knowledge: distributed denial-of-service (DDoS) attack [M2]**
    - Developed a categorizing schema by systematically selecting and examining ∼24,000 works related to volumetric DDoS attack, derived from DBLP and Google Scholar.
    - Performed detailed analysis over selected to derive systematized repeating design patterns and a set of IP network architecture properties.
    - Categorized the above papers into sub-categories based on their deployment locations, approaches, incentives, etc. and contributed to writing a research paper.
  - **Scaling Transport-Layer protocol in Named Data Network (NDN) [P4]**
    - Designed and implemented both randomized and most recent partial-states States Vector Sync (p-SVS) to scale with a large number of data producers within the same group.
    - Simulated experiments on p-SVS over an NDN simulation tool named ndnSIM over several topologies.

Industry Experience

- **Arista Networks, Inc.** — Los Angeles, CA, USA
  Software Engineer Intern — *Jun. 2022 - Sep.2022*

  IEEE 802.1Q Tunneling CLI
  - Designed the new module architecture that significantly reduced the code complexity compared to the existing similar tunneling implementation and completed a detailed design document.
  - Implemented software-side reactors and hardware-side bit setter that together can filter packets violating user-defined VLAN rules in 802.1Q tunneling. (∼10k lines of codes)
  - Pushed the changes to the next release to be used by all switches over a specific popular platform.

- **Taboola, Inc.** — Los Angeles, CA, USA
  Machine Learning/Data Science Intern — *Jun. 2019 - Sep. 2019*

  Knowledge Base of News Keywords
  - Built an end-to-end pipeline with Spark SQL and Java to process data crawled by IBM Watson and further capture their embeddings with Word2Vec. (∼5k lines of codes)
  - Devised algorithms for de-duplicating keywords based on a combined metric, including similar neighbors, lexical similarity, etc.
  - Proposed a Knowledge Base representation of news keywords over Neo4j to effectively visualize keywords relationships and implemented an auto-renewal process that runs daily.

- **Qihoo 360 Technology Co.** — Beijing, China
  Machine Learning/Data Science Intern — *Jun. 2018 - Sep. 2018*

  Internet Traffic Classification and Anomaly Detection
  - Conducted surveys, implementations, and experiments on state-of-art machine learning algorithms for traffic anomaly detection and manually examined benign and malicious internet traffic samples.
  - Selected features and devised an n-grams algorithm to form pseudo images from traffic.
  - Designed a Random Forest model and a Neural Network model to achieve a 4% false positive rate and a 94% true positive rate.

Teaching Experience

- **COM SCI 35L: Software Construction Laboratory** — Los Angeles, CA, USA
  Teaching Assistant — *Fall 2021*
  - Lectured 20 hours of material focusing on Git, Shell, Vim, Java, etc., to 52 students and held 20 hours of office hours for ∼250 students.
  - Mentored ∼10 groups of undergraduate students completing Node.js/React projects.
  - Graded ∼250 students' coding assignments and 2 exams.

Misc

- **Selected Courses**: Database System, Introduction to Machine Learning, Operating Systems, Compiler Construction, Internet Architecture and Protocols, Current Topics in Computer System Modeling Analysis.
- **Selected Languages**: Python, C/C++, Java, SQL, Bash, Datalog.
- **Selected Platforms**: Amazon EC2, Sklearn, Github, Neo4j, Apache Spark, Apache Flink, Spark Streaming, LaTeX.